

Mark Valinh Phouiphiphath

mark21097@gmail.com • (805)340-8259 • www.linkedin.com/in/mark-phouiphiphath • github.com/mark21097

EDUCATION

University of California, Irvine

March 2027

B.S. in Computer Science

Relevant Courses: Machine Learning, AI, Linear Algebra, Statistics, Database Management, Software Engineering, Data Structures

TECHNICAL SKILLS

Languages: C/C++, Java, HTML/CSS/JavaScript, **Python**, R, **SQL**

Data Science & ML: NumPy, Pandas, **Scikit-learn**, TensorFlow, **PyTorch**, NLP, **Neural Networks**, **XGBoost**, MLflow

Web & Systems: React, Node.js/Fastify, **FastAPI**, Firebase, **AWS S3**, **PostgreSQL**

Tools: Git/GitHub/GitLab, **Jupyter Notebook**, Linux, **Docker**

EXPERIENCE

Boundary RSS | Remote

January 2026 - Present

Data Scientist & ML Engineer Intern

- Analyzed large-scale structured sensor and operational datasets to support research and development initiatives, improving data reliability and downstream modeling performance by **30%**.
- Built, trained, and evaluated classification and regression models using Python, Pandas, NumPy, and scikit-learn, achieving up to **25% improvement** in predictive accuracy over baseline models.
- Designed data preprocessing and feature engineering pipelines (outlier handling, normalization, feature selection) that reduced noise and improved model generalization by **20%**.
- Collaborated within an **Agile** sprint framework, including weekly standups, sprint planning, and iterative reviews to deliver data pipeline milestones on schedule across cross-functional engineering teams.

Fusion ICS Project | Irvine, CA

Nov 2024 - May 2025

Project Developer – Full-Stack Development & Data-Driven Applications

- Built and deployed a full-stack mentorship platform for **100+ students**, integrating Firebase Firestore for real-time data sync, role-based authentication, and a swipe-to-match algorithm filtering by role, interest, and availability.
- Translated **15+ Figma wireframes** into responsive UI components, enforcing reusable design patterns that reduced feature development time by **25%** and maintained visual consistency across all core application views.
- Collaborated with a team of **10+ developers/designers** in an Agile workflow to ship profile creation with resume integration, location/role filtering, and a full matching pipeline end-to-end.

PROJECTS

ChurnSight | Python, Scikit-learn, PostgreSQL, FastAPI, Docker

January 2026

- Built an end-to-end churn prediction pipeline on 50K+ subscriber records stored in PostgreSQL, performing EDA, preprocessing, and SMOTE-based class balancing to ensure analytics-ready inputs, improving minority class recall by **32%**.
- Benchmarked Random Forest, XGBoost, and Logistic Regression models tracked via **MLflow**, achieving **91% AUC-ROC** through GridSearchCV tuning and cross-validation, reducing iteration cycles by **25%**.
- Containerized the full training and inference pipeline with **Docker** and deployed a **FastAPI** REST endpoint for real-time churn scoring, enabling scalable, production-ready predictions and cutting environment setup time by **40%**.

TalentLens | Python, Pandas, SQL

October 2025

- Built an end-to-end data science pipeline analyzing **100,000+ job postings**, performing EDA to identify hiring trends, skill demand patterns, and salary distributions across software and data roles.
- Performed extensive data cleaning and feature engineering on unstructured text (job descriptions), extracting skill keywords and role metadata, improving signal quality by **40%**.
- Trained supervised NLP-based classification models to predict job role category (SWE, Data Scientist, ML Engineer) with **85%+ accuracy** using cross-validation and regularization tuning.